

Programming GPUs with PyCuda

Nicolas Pinto (MIT) and Andreas Klöckner (Brown)

SciPy Conference 2009 / Advanced Tutorial
http://conference.scipy.org/advanced_tutorials

August 19, 2009

Thanks

- SciPy2009 Organizers
- **Andreas Klöckner (!)**
- PyCuda contributors
- Nvidia Corporation

Outline

- 1 Introduction
- 2 Programming GPUs
- 3 GPU Scripting
- 4 PyCuda Hands-on: Matrix Multiplication

Outline

- 1 Introduction
 - GPU Computing: Overview
 - Architectures and Programming Models
- 2 Programming GPUs
- 3 GPU Scripting
- 4 PyCuda Hands-on: Matrix Multiplication

Outline

1 Introduction

- GPU Computing: Overview
- Architectures and Programming Models


2 Programming GPUs

3 GPU Scripting


4 PyCuda Hands-on: Matrix Multiplication

Stream Processing?

- Design target for CPUs:
 - Focus on *Task* parallelism
 - Make a single thread very fast
 - Hide latency through large caches
 - Predict, speculate




CPU




GPU

Stream Processing?

- Design target for CPUs:
 - Focus on *Task* parallelism
 - Make a single thread very fast
 - Hide latency through large caches
 - Predict, speculate
- *Stream* Processing takes a different approach:
 - Focus on *Data* parallelism
 - Throughput matters – single threads do not
 - Hide latency through parallelism
 - Let programmer deal with “raw” memory hierarchy




CPU



GPU


CPU Chip Real Estate



Die floorplan: Intel Core i7 (2008).

45 nm, 4x4 SP ops at a time, 4x256KB L2, 8MB L3

GPU Chip Real Estate



Die floorplan: AMD RV770 (2008).
55 nm, 800 SP ops at a time.

Market Overview

Quote Linus Torvalds:

“Hardware that isn’t mass market tends to not be worth it in the long run.”

Market Overview

Quote Linus Torvalds:

“Hardware that isn’t mass market tends to not be worth it in the long run.”

Based on that:

- Sony/Toshiba/IBM: Cell Broadband Engine
- ATI: R580 and later
- Nvidia: G80 and later
- Intel: Larabee

Outline

1 Introduction


- GPU Computing: Overview
- Architectures and Programming Models

2 Programming GPUs

3 GPU Scripting

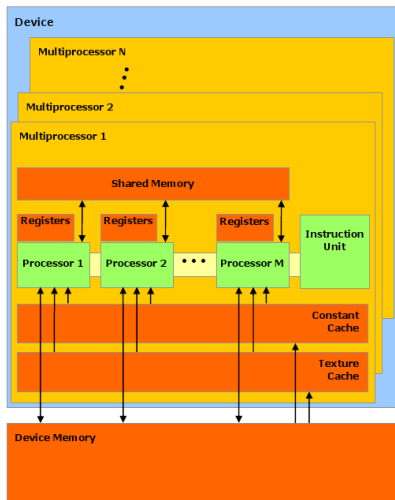
4 PyCuda Hands-on: Matrix Multiplication

Cell BE: Architecture




- 1 Cell BE = 1 dual-core Power + 8 SPEs + Bus
- 1 SPE = SPU + DMA + 256 KiB Local Store
- 1 SPU = 128-bit Vector ALU
- Bus = 200 GB/s Ring
- Ded. RAM (25 GB/s)

GPU: Architecture (e.g. Nvidia)



- 1 GPU = 30 MPs
- 1 MP = 1 ID (1/4 clock) + 8 SP + 1 DP + 16 KiB Shared + 32 KiB Reg + HW Sched
- Scalar cores
max 512 threads/MP
- Ded. RAM (140 GB/s)
- PCIe2 Host DMA (6 GB/s)
- Limited Caches

Intel Larabee: Architecture



- Unreleased (2010?)
- x86-64 + SSE +
“vector-complete” 512-bit ISA (“LRBni”)
- 4x “Hyperthreading”
- 32 (?) cores per chip
- “Fiber/Strand” software threads
- Recursive Launches
- Coherent Caches (w/ explicit control)
- Performance?

Programming Models



Pixel Shaders?




DirectX?




OpenGL?

Programming Models




Programming Models




- Dedicated Compute APIs
- Not much “graphicsy” stuff visible

Programming Models



Hardware-specific ←————→ Abstract

Programming Models



Architecture Comparison

Cell	GPU	Larabee
<ul style="list-style-type: none">⊕ ~ Multicore⊕ Open Spec⊖ Hard: DMA sched, Alignment, Small LS⊖ HW Avail. (\$)⊖ Mem BW		

Architecture Comparison

Cell	GPU	Larabee
<ul style="list-style-type: none"> ⊕ ~ Multicore ⊕ Open Spec ⊖ Hard: DMA sched, Alignment, Small LS ⊖ HW Avail. (\$) ⊖ Mem BW 	<ul style="list-style-type: none"> ⊕ Available now ⊕ OpenCL ⊕ Maturing (tools etc.) ⊕ Mem BW ⊙ Reasonably easy ⊖ Perf. opaque 	

Architecture Comparison

Cell	GPU	Larabee
<ul style="list-style-type: none"> ⊕ ~ Multicore ⊕ Open Spec ⊖ Hard: DMA sched, Alignment, Small LS ⊖ HW Avail. (\$) ⊖ Mem BW 	<ul style="list-style-type: none"> ⊕ Available now ⊕ OpenCL ⊕ Maturing (tools etc.) ⊕ Mem BW ⊕ Reasonably easy ⊖ Perf. opaque 	<ul style="list-style-type: none"> ⊖ Unavailable ⊕ Programmability? ⊕ OpenCL Support (→ why wait?) ⊖ Competitive with next-gen GPU? (→ why specialize?)

Architecture Comparison

Cell	GPU	Larabee
<ul style="list-style-type: none"> ⊕ ~ Multicore ⊕ Open Spec ⊖ Hard: DMA sched, Alignment, Small LS ⊖ HW Avail. (\$) ⊖ Mem BW 	<ul style="list-style-type: none"> ⊕ Available now ⊕ OpenCL ⊕ Maturing (tools etc.) ⊕ Mem BW ⊕ Reasonably easy ⊖ Perf. opaque <div style="border: 2px solid red; padding: 5px; display: inline-block; transform: rotate(-2deg); color: red; font-weight: bold;">Focus Here</div>	<ul style="list-style-type: none"> ⊖ Unavailable ⊕ Programmability? ⊕ OpenCL Support (→ why wait?) ⊖ Competitive with next-gen GPU? (→ why specialize?)

Questions?

?

Outline

1 Introduction

2 Programming GPUs

- Overview
- Dealing with Memory

3 GPU Scripting

4 PyCuda Hands-on: Matrix Multiplication

Outline

1 Introduction

2 Programming GPUs

■ Overview

■ Dealing with Memory

3 GPU Scripting

4 PyCuda Hands-on: Matrix Multiplication

What is CUDA?

- CUDA is Nvidia's proprietary compute abstraction.

What is CUDA?

- CUDA is Nvidia's proprietary compute abstraction.
- Main merit: A well-balanced model of GPU computing.
 - Abstract enough to not be hardware-specific.
 - Concrete enough to expose most hardware features.

What is CUDA?

- CUDA is Nvidia's proprietary compute abstraction.
- Main merit: A well-balanced model of GPU computing.
 - Abstract enough to not be hardware-specific.
 - Concrete enough to expose most hardware features.
- (Very) close semantic relative of OpenCL.

Gains and Losses

Gains

- ➕ Memory Bandwidth
(140 GB/s vs. 12 GB/s)
- ➕ Compute Bandwidth
(Peak: 1 TF/s vs. 50 GF/s,
Real: 200 GF/s vs. 10 GF/s)

Losses

Gains and Losses

Gains

- ➕ Memory Bandwidth
(140 GB/s vs. 12 GB/s)
- ➕ Compute Bandwidth
(Peak: 1 TF/s vs. 50 GF/s,
Real: 200 GF/s vs. 10 GF/s)


Losses

- ➖ Recursion
- ➖ Function pointers
- ➖ Exceptions
- ➖ IEEE 754 FP compliance
- ➖ Cheap branches (i.e. ifs)

GPUs: Threading Model


■ Multi-tiered Parallelism

- *Grid* of Blocks
- *Block* of Threads




GPUs: Threading Model

- Multi-tiered Parallelism
 - *Grid* of Blocks
 - *Block* of Threads
- Only threads within a block can communicate




GPUs: Threading Model

- Multi-tiered Parallelism
 - *Grid* of Blocks
 - *Block* of Threads
- Only threads within a block can communicate
 - Each Block is assigned to a physical execution unit.




GPUs: Threading Model

- Multi-tiered Parallelism
 - *Grid* of Blocks
 - *Block* of Threads
- Only threads within a block can communicate
 - Each Block is assigned to a physical execution unit.
- Algorithm must work with blocks executed in any order




GPUs: Threading Model



- Multi-tiered Parallelism
 - *Grid* of Blocks
 - *Block* of Threads
- Only threads within a block can communicate
 - Each Block is assigned to a physical execution unit.
- Algorithm must work with blocks executed in any order
- Grids and Blocks replace outer loops in an algorithm.

GPUs: Threading Model



- Multi-tiered Parallelism
 - *Grid* of Blocks
 - *Block* of Threads
- Only threads within a block can communicate
 - Each Block is assigned to a physical execution unit.
- Algorithm must work with blocks executed in any order
- Grids and Blocks replace outer loops in an algorithm.
- Indices available at run time

My first CUDA program

```
1 // GPU-side
2
3 __global__ void square_array(float *a, int n)
4 {
5     int i = blockIdx.x * blockDim.x + threadIdx.x;
6     if (i < n)
7         a[i] = a[i] * a[i];
8 }
```

My first CUDA program

```
12 int main() // CPU-side
13 {
14     cudaSetDevice(0); // EDIT ME
15
16     int n = 4096; int bytes = n*sizeof( float );
17     float *a_host = ( float *) malloc(bytes);
18     for ( int i = 0; i < n; i++) a_host[i] = i;
19
20     float *a_device;
21     cudaMalloc((void **) &a_device, bytes);
22     cudaMemcpy(a_device, a_host, bytes, cudaMemcpyHostToDevice);
23
24     int block_size = 256;
25     int n_blocks = (n + block_size-1) / block_size;
26     square_array<<<n_blocks, block_size>>>(a_device, n);
27
28     free(a_host); cudaFree(a_device);
29 }
```


A Dose of Reality: Error Checking

```
#define CUDA_CHK(NAME, ARGS) { \
    cudaError_t cuda_err_code = NAME ARGS; \
    if (cuda_err_code != cudaSuccess) { \
        printf ("%s failed with code %d\n", #NAME, cuda_err_code); \
        abort (); \
    } \
}

CUDA_CHK(cudaMalloc, (&result, m_size*sizeof(float)));
```

Typical errors:

- GPUs have (some) memory protection → “Launch failure”
- Invalid sizes (block/grid/...)

Invisible Subtleties

Host Pointer or Device Pointer?

```
float *h_data = (float*) malloc(mem_size);  
float *d_data;  
CUDA_CHK(cudaMalloc, ((void**) &d_data, mem_size));  
→ Both kinds of pointer share the same data type!
```

Invisible Subtleties

Host Pointer or Device Pointer?

```
float *h_data = (float*) malloc(mem_size);  
float *d_data;  
CUDA_CHK(cudaMalloc, ((void**) &d_data, mem_size));  
→ Both kinds of pointer share the same data type!
```

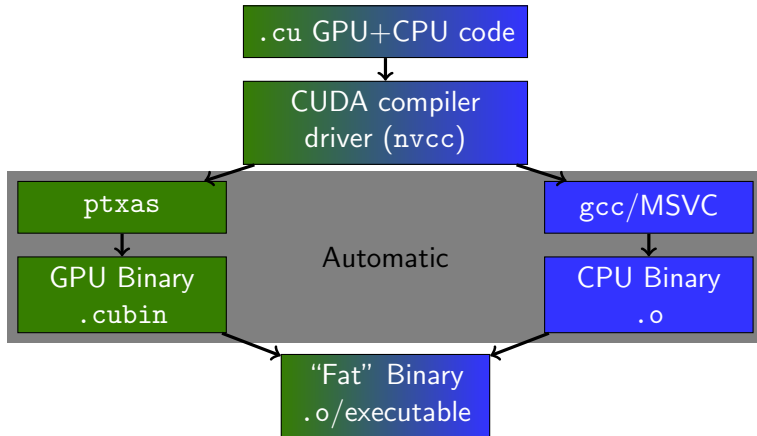
Kernel Launches

Execution configuration:


```
dim3 grid_size(gx, gy); // max 2D  
dim3 block_size(bx, by, bz); // max 3D  
kernel <<<grid_size, block_size>>>(arg, ...);
```

- Do not wait for completion.
- Cheap! ($\sim 2 \mu\text{s}$ overhead)

The CUDA Toolchain



Executing CUDA Binaries



GPU Demo Machines

Machine	GPUs
iapcuda-01	Device 0: " GeForce GTX 285"
iapcuda-01	Device 1: " Tesla C1060"
iapcuda-01	Device 2: " Tesla C1060"
iapcuda-02	Device 0: " GeForce GTX 295"
iapcuda-02	Device 1: " GeForce GTX 295"
iapcuda-02	Device 2: " Tesla C1060"
iapcuda-02	Device 3: " Tesla C1060"

Prepare your workspace in one of our CUDA demo machine:

```
1 ssh scipy09@iapcuda-NN.no-ip.org  
    (password: GpUh4cK3r)
```

GPU Demo Machines

Machine	GPUs
iapcuda-01	Device 0: " GeForce GTX 285"
iapcuda-01	Device 1: " Tesla C1060"
iapcuda-01	Device 2: " Tesla C1060"
iapcuda-02	Device 0: " GeForce GTX 295"
iapcuda-02	Device 1: " GeForce GTX 295"
iapcuda-02	Device 2: " Tesla C1060"
iapcuda-02	Device 3: " Tesla C1060"

Prepare your workspace in one of our CUDA demo machine:

- 1 `ssh scipy09@iapcuda-NN.no-ip.org`
(password: GpUh4cK3r)
- 2 `mkdir lastname.firstname`

GPU Demo Machines

Machine	GPUs
iapcuda-01	Device 0: " GeForce GTX 285"
iapcuda-01	Device 1: " Tesla C1060"
iapcuda-01	Device 2: " Tesla C1060"
iapcuda-02	Device 0: " GeForce GTX 295"
iapcuda-02	Device 1: " GeForce GTX 295"
iapcuda-02	Device 2: " Tesla C1060"
iapcuda-02	Device 3: " Tesla C1060"

Prepare your workspace in one of our CUDA demo machine:

- 1 `ssh scipy09@iapcuda-NN.no-ip.org`
(password: GpUh4cK3r)
- 2 `mkdir lastname.firstname`
- 3 `cd lastname.firstname`

GPU Demo Machines

Machine	GPUs
iapcuda-01	Device 0: " GeForce GTX 285"
iapcuda-01	Device 1: " Tesla C1060"
iapcuda-01	Device 2: " Tesla C1060"
iapcuda-02	Device 0: " GeForce GTX 295"
iapcuda-02	Device 1: " GeForce GTX 295"
iapcuda-02	Device 2: " Tesla C1060"
iapcuda-02	Device 3: " Tesla C1060"

Prepare your workspace in one of our CUDA demo machine:

- 1 `ssh scipy09@iapcuda-NN.no-ip.org`
(password: GpUh4cK3r)
- 2 `mkdir lastname.firstname`
- 3 `cd lastname.firstname`
- 4 `wget is.gd/2o40o && tar xzf scipy09-pycuda-tut.tar.gz`

Getting your feet wet

Hands-on Exercise

- 1 Edit 1-cuda-simple/simple.cu:
`cudaSetDevice(Your GPU #);`
- 2 Compile and run:
`nvcc -o simple.x simple.cu`
`./simple.x`
- 3 Add error checking to the example.
- 4 Modify simple.cu to print the contents of the result.
- 5 Modify simple.cu to compute $c_i = a_i b_i$.
- 6 Modify simple.cu to use blocks of 16×16 threads.

My first CUDA program (Solution)

Preliminary bits:

```
1 #include <stdio.h>
2
3 #define CUDA_CHK(NAME, ARGS) { \
4     cudaError_t cuda_err_code = NAME ARGS; \
5     if (cuda_err_code != cudaSuccess) { \
6         printf("%s failed with code %d\n", #NAME, cuda_err_code); \
7         abort(); \
8     } \
9 }
```

My first CUDA program (Solution)

The GPU kernel:

```
13 __global__ void square_array (float *a, float *b, int n)
14 {
15     int i = (blockIdx.x * blockDim.y + threadIdx.y)
16         * blockDim.x + threadIdx.x;
17     if (i < n)
18         a[i] = a[i] * b[i];
19 }
```

My first CUDA program (Solution)

Allocating memory:

```
23 int main()
24 {
25     cudaSetDevice(0); // EDIT ME
26
27     const int n = 4096;
28
29     float *a_host = (float *) malloc(n*sizeof(float ));
30     float *b_host = (float *) malloc(n*sizeof(float ));
31
32     float *a_device, *b_device;
33     CUDA_CHK(cudaMalloc, ((void **) &a_device, n*sizeof(float)));
34     CUDA_CHK(cudaMalloc, ((void **) &b_device, n*sizeof(float)));
```

My first CUDA program (Solution)

Transfer and Launch:

```
38  for (int i = 0; i < n; i++) { a_host[i] = i; b_host[i] = i+1; }
39
40  CUDA_CHK(cudaMemcpy, (a_device, a_host, n*sizeof(float),
41      cudaMemcpyHostToDevice));
42  CUDA_CHK(cudaMemcpy, (b_device, b_host, n*sizeof(float),
43      cudaMemcpyHostToDevice));
44
45  dim3 block_dim(16, 16);
46  int block_size = block_dim.x*block_dim.y;
47  int n_blocks = (n + block_size-1) / block_size;
48  square_array<<<n_blocks, block_dim>>>(a_device, b_device, n);
```

My first CUDA program (Solution)

Output and Clean-up:

```
52  CUDA_CHK(cudaMemcpy, (a_host, a_device, n*sizeof(float),
53      cudaMemcpyDeviceToHost));
54
55  for (int i = 0; i < n; i++)
56      printf("%.0f ", a_host[i]);
57  puts("\n");
58
59  free(a_host);
60  CUDA_CHK(cudaFree, (a_device));
61 }
```

Questions?

?

Outline

1 Introduction

2 Programming GPUs

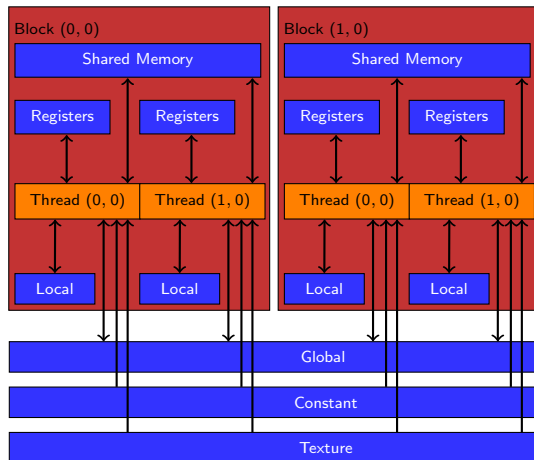
■ Overview

■ Dealing with Memory

3 GPU Scripting

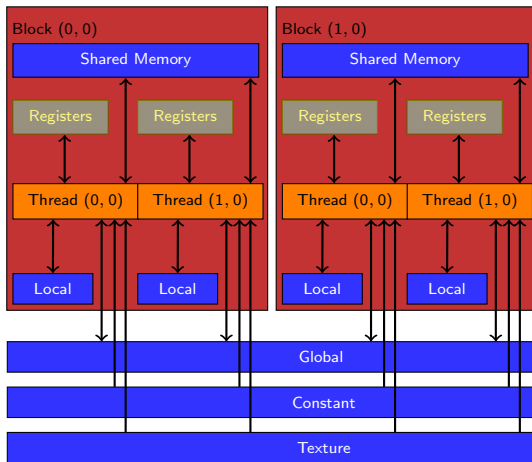
4 PyCuda Hands-on: Matrix Multiplication

Memory Model



Already seen:

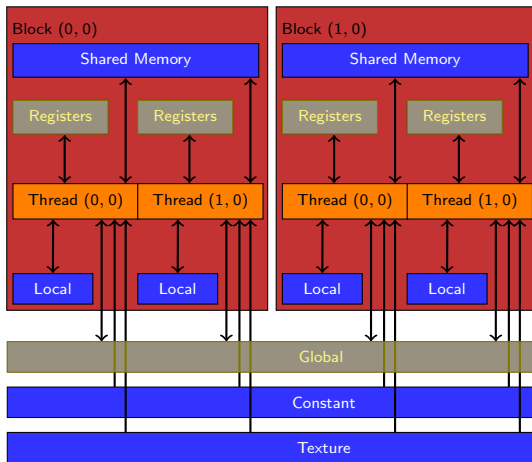
Memory Model



Already seen:

- Registers

Memory Model

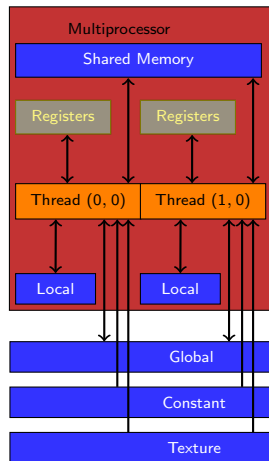


Already seen:

- Registers
- Global

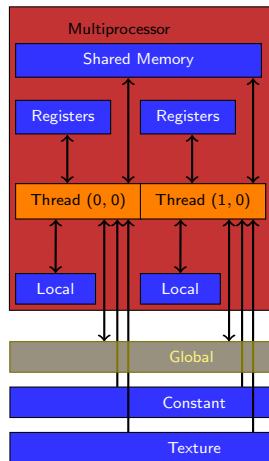
Registers

- 32 KiB of registers per MP
- Per-thread
- Latency: 1 clock
- Variable amount per thread
 - Register count limits max. threads/MP
 - CPUs: Fixed register file (\sim)




Global Memory

- Several GiB usually
- Per-GPU
- Latency: ~ 1000 clocks
- 512 bit memory bus
 - Best throughput: 16 consecutive threads read aligned chunk



Example: Matrix Transpose



Naive: Using global memory

First attempt: Naive port of the CPU code.

```
__global__ void transpose(float *out, float *in, int w, int h) {  
    unsigned int xldx = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yldx = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if ( xldx < w && yldx < h ) {  
        unsigned int idx_in  = xldx + w * yldx;  
        unsigned int idx_out = yldx + h * xldx;  
  
        out[idx_out] = in[idx_in ];  
    }  
}
```


Measuring Performance

Writing high-performance Codes

Mindset: What is going to be the limiting factor?

- Floating point throughput?
- Memory bandwidth?
 - Cache sizes?

Measuring Performance

Writing high-performance Codes

Mindset: What is going to be the limiting factor?

- Floating point throughput?
- Memory bandwidth?
 - Cache sizes?

Benchmark the assumed limiting factor right away.

Measuring Performance

Writing high-performance Codes

Mindset: What is going to be the limiting factor?

- Floating point throughput?
- Memory bandwidth?
 - Cache sizes?


Benchmark the assumed limiting factor right away.

Evaluate

- Know your peak throughputs (roughly)
- Are you getting close?
- Are you tracking the right limiting factor?

Performance: Matrix transpose

Very likely: Bound by memory bandwidth.



Fantastic! About same as CPU. Why?

Naive: Using global memory

```
__global__ void transpose(float *out, float *in, int w, int h) {  
    unsigned int xldx = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yldx = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if ( xldx < w && yldx < h ) {  
        unsigned int idx_in  = xldx + w * yldx;  
        unsigned int idx_out = yldx + h * xldx;  
  
        out[idx_out] = in[idx_in];  
    }  
}
```

Naive: Using global memory

```
__global__ void transpose(float *out, float *in, int w, int h) {  
    unsigned int xldx = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yldx = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if ( xldx < w && yldx < h ) {  
        unsigned int idx_in  = xldx + w * yldx;  
        unsigned int idx_out = yldx + h * xldx;  
  
        out[idx_out] = in[idx_in];  
    }  
}
```

Reading from global mem:



stride: 1

Naive: Using global memory

```
__global__ void transpose(float *out, float *in, int w, int h) {  
    unsigned int xldx = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yldx = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if ( xldx < w && yldx < h ) {  
        unsigned int idx_in  = xldx + w * yldx;  
        unsigned int idx_out = yldx + h * xldx;  
  
        out[idx_out] = in[idx_in];  
    }  
}
```

Reading from global mem:



stride: 1 \rightarrow one mem.trans.

Naive: Using global memory


```
__global__ void transpose(float *out, float *in, int w, int h) {  
    unsigned int xldx = blockDim.x * blockIdx.x + threadIdx.x;  
    unsigned int yldx = blockDim.y * blockIdx.y + threadIdx.y;  
  
    if ( xldx < w && yldx < h ) {  
        unsigned int idx_in  = xldx + w * yldx;  
        unsigned int idx_out = yldx + h * xldx;  
  
        out[idx_out] = in[idx_in];  
    }  
}
```

Reading from global mem:



stride: 1 \rightarrow one mem.trans.

Writing to global mem:



stride: 16

Naive: Using global memory

```
__global__ void transpose(float *out, float *in, int w, int h) {
    unsigned int xldx = blockDim.x * blockIdx.x + threadIdx.x;
    unsigned int yldx = blockDim.y * blockIdx.y + threadIdx.y;

    if ( xldx < w && yldx < h ) {
        unsigned int idx_in  = xldx + w * yldx;
        unsigned int idx_out = yldx + h * xldx;


        out[idx_out] = in[idx_in ];
    }
}
```

Reading from global mem:



stride: 1 \rightarrow one mem.trans.

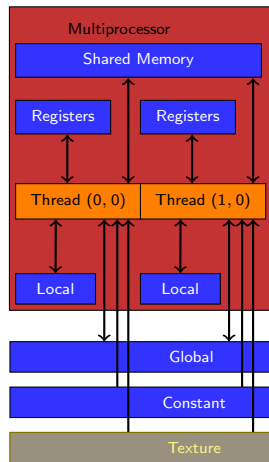
Writing to global mem:



stride: 16 \rightarrow **16 mem.trans.!**

Texture Memory

- Same memory as global
- But: more access patterns achieve usable bandwidth
- Optional: 2D and 3D indexing
- Small, incoherent Cache (prefers n D-local access)
- Read-only
- Latency: ~ 1000 clocks (despite cache!)
- Optional: Linear Interpolation



Transpose with Textures

```
texture <float, 1, cudaReadModeElementType> in_tex;

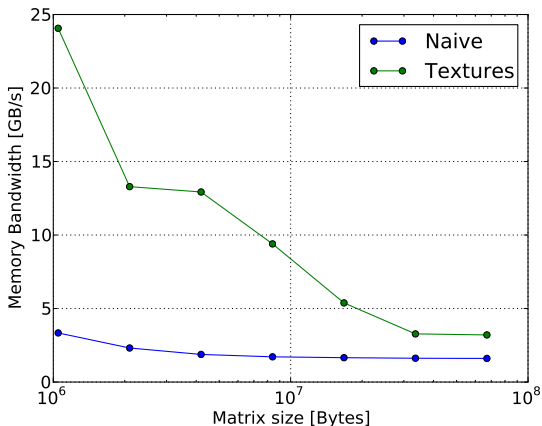
__global__ void transpose( float *out, int w, int h ) {
    unsigned int yldx = blockDim.x * blockIdx.x + threadIdx.x;
    unsigned int xldx = blockDim.y * blockIdx.y + threadIdx.y;

    if ( xldx < w && yldx < h ) {
        unsigned int idx_in  = xldx + w * yldx;
        unsigned int idx_out = yldx + h * xldx;

        out[idx_out] = tex1Dfetch(in_tex, idx_in );
    }
}

#define PREPARE \
    cudaBindTexture(0, in_tex, d_idata, mem_size); \
    std::swap(grid.x, grid.y); \
    std::swap(threads.x, threads.y);
```

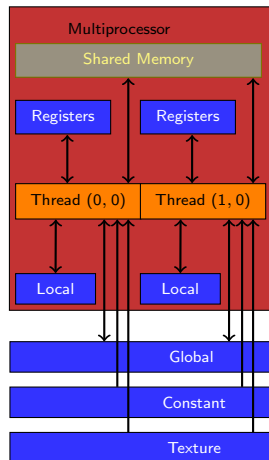
Performance: Transpose with Textures



Better! But texture units can't quite hide wide data bus.
Need different idea.

Shared Memory

- 16 KiB of shared mem per MP
- Per-block
- Latency: 2 clocks
- Variable amount per block
 - Shared memory limits max. blocks/MP
- Banked



Transpose: Idea


- Global memory dislikes non-unit strides.
- Shared memory doesn't mind.

Idea

- Don't transpose element-by-element.
- Transpose block-by-block instead.

- 1 Read untransposed block from global and write to shared
- 2 Read block transposed from shared and write to global

Illustration: Blockwise Transpose



Improved: Using shared memory

```
__global__ void transpose( float *out, float *in, int w, int h ) {
    __shared__ float block[BLOCK_DIM*BLOCK_DIM];

    unsigned int xBlock = blockDim.x * blockIdx.x;
    unsigned int yBlock = blockDim.y * blockIdx.y;

    unsigned int xIndex = xBlock + threadIdx.x;
    unsigned int yIndex = yBlock + threadIdx.y;


    unsigned int index_out, index_transpose;

    if ( xIndex < w && yIndex < h ) {
        unsigned int index_in = w * yIndex + xIndex;
        unsigned int index_block = threadIdx.y * BLOCK_DIM + threadIdx.x;

        block[index_block] = in[index_in];
        index_transpose = threadIdx.x * BLOCK_DIM + threadIdx.y;
        index_out = h * (xBlock + threadIdx.y) + yBlock + threadIdx.x;
    }
    __syncthreads();

    if ( xIndex < w && yIndex < h ) {
        out[index_out] = block[index_transpose];
    }
}
```


Performance: Transpose with Shared Memory



Not bad! Are we done?

Review: Memory Model

Type	Per	Access	Latency	
Registers	thread	R/W	1	
Local	thread	R/W	1000	
Shared	block	R/W	2	
Global	grid	R/W	1000	Not cached
Constant	grid	R/O	1-1000	Cached
Texture	grid	R/O	1000	Spatially cached

Important

Don't “*choose one*” type of memory.

Successful algorithms combine many types' strengths.

Questions?

?

Outline

1 Introduction

2 Programming GPUs

3 GPU Scripting

- Scripting + GPUs: A good combination
- Whetting your Appetite
- Working with PyCuda
- A peek under the hood
- Metaprogramming CUDA

4 PyCuda Hands-on: Matrix Multiplication

Outline

- 1 Introduction
- 2 Programming GPUs
- 3 GPU Scripting
 - Scripting + GPUs: A good combination
 - Whetting your Appetite
 - Working with PyCuda
 - A peek under the hood
 - Metaprogramming CUDA
- 4 PyCuda Hands-on: Matrix Multiplication

Scripting Languages

Python:

- is discoverable and interactive.
- has comprehensive built-in functionality.
- manages resources automatically.
- uses run-time typing.
- works well for “gluing” lower-level blocks together.

Scripting: Goals

Scripting languages aim to reduce the load on the programmer:

- Reduce required knowledge
- Encourage experimentation
- Eliminate sources of error
- Encourage abstraction wherever possible
- Value programmer time over computer time

Think about the tools you use.

Use the right tool for the job.

Our mantra: always use the right tool !



Why do Scripting for GPUs?

- GPUs are everything that scripting languages are not.
 - Highly parallel
 - Very architecture-sensitive
 - Built for maximum compute/memory throughput→ complement each other
- CPU: largely restricted to control tasks ($\sim 1000/\text{sec}$)
 - Scripting fast enough
- Realize a promise: Use Scripting...
 - from first prototype
 - to full-scale production code.



Scripting: Speed

- Usual answer to the “Speed Question”:
Hybrid (“mixed”) Code.
- Plays to the strengths of each language.
- But: Introduces (some) complexity.



Observation: GPU code is already hybrid.

Consequence: No added complexity through hybrid code.

Questions?

?

Outline

1 Introduction

2 Programming GPUs

3 GPU Scripting

- Scripting + GPUs: A good combination
- **Whetting your Appetite**
- Working with PyCuda
- A peek under the hood
- Metaprogramming CUDA

4 PyCuda Hands-on: Matrix Multiplication

Whetting your appetite

```
1 import pycuda.driver as cuda
2 import pycuda.autoinit
3 import numpy
4
5 a = numpy.random.randn(4,4).astype(numpy.float32)
6 a_gpu = cuda.mem_alloc(a.nbytes)
7 cuda.memcpy_htod(a_gpu, a)
```

[This is `examples/demo.py` in the PyCuda distribution.]

Whetting your appetite

```
1 mod = cuda.SourceModule("""
2     __global__ void doublify ( float *a)           Compute kernel
3     {
4         int idx = threadIdx.x + threadIdx.y*4;
5         a[idx] *= 2;
6     }
7     """)
8
9 func = mod.get_function("doublify")
10 func(a_gpu, block=(4,4,1))
11
12 a_doubled = numpy.empty_like(a)
13 cuda.memcpy_dtoh(a_doubled, a_gpu)
14 print a_doubled
15 print a
```

Whetting your appetite, Part II

Did somebody say “Abstraction is good”?

Whetting your appetite, Part II

```
1 import numpy
2 import pycuda.autoinit
3 from pycuda import gpuarray
4
5 a_cpu = numpy.random.randn(4,4).astype(numpy.float32)
6 b_cpu = numpy.random.randn(4,4).astype(numpy.float32)
7 c_cpu = a_cpu * b_cpu
8
9 a_gpu = gpuarray.to_gpu(a_cpu)
10 b_gpu = gpuarray.to_gpu(b_cpu)
11 c_gpu = (a_gpu * b_gpu).get()
12
13 print c_cpu - c_gpu
```


Remember me?

```

1 // trivia
2 #include <stdio.h>
3
4 #define CUDA_CHK(NAME, ARGS) { \
5     cudaError_t cuda_err_code = NAME ARGS; \
6     if (cuda_err_code != cudaSuccess) { \
7         printf("%s failed with code %d\n", #NAME, cuda_err_code); \
8         abort(); \
9     } \
10 } \
11 // end
12
13 // kernel
14 _global_ void square_array(float *a, float *b, int n)
15 {
16     int i = (blockIdx.x * blockDim.y + threadIdx.y)
17         * blockDim.x + threadIdx.x;
18     if (i < n)
19         a[i] = a[i] * b[i];
20 }
21 // end
22
23 // main1
24 int main()
25 {
26     cudaSetDevice(0); // EDIT ME
27
28     const int n = 4096;
29
30     float *a_host = (float *) malloc(n*sizeof(float));
31     float *b_host = (float *) malloc(n*sizeof(float));
32
33     float *a_device, *b_device;
34     CUDA_CHK(cudaMalloc, ((void **) &a_device, n*sizeof(float)));
35     CUDA_CHK(cudaMalloc, ((void **) &b_device, n*sizeof(float)));
36 // end

```

```

1 // main2
2 for (int i = 0; i < n; i++) { a_host[i] = i; b_host[i] = i+1; }
3
4 CUDA_CHK(cudaMemcpy, (a_device, a_host, n*sizeof(float),
5     cudaMemcpyHostToDevice));
6 CUDA_CHK(cudaMemcpy, (b_device, b_host, n*sizeof(float),
7     cudaMemcpyHostToDevice));
8
9 dim3 block_dim(16, 16);
10 int block_size = block_dim.x*block_dim.y;
11 int n_blocks = (n + block_size - 1) / block_size;
12 square_array <<<n_blocks, block_dim>>>(a_device, b_device, n);
13 // end
14
15 // main3
16 CUDA_CHK(cudaMemcpy, (a_host, a_device, n*sizeof(float),
17     cudaMemcpyDeviceToHost));
18
19 for (int i = 0; i < n; i++)
20     printf("%0f ", a_host[i]);
21 puts("\n");
22
23 free(a_host);
24 CUDA_CHK(cudaFree, (a_device));
25 }
26 // end

```

Outline

1 Introduction


2 Programming GPUs

3 GPU Scripting

- Scripting + GPUs: A good combination
- Whetting your Appetite
- **Working with PyCuda**
- A peek under the hood
- Metaprogramming CUDA

4 PyCuda Hands-on: Matrix Multiplication

PyCuda Philosophy



- Provide complete access
- Automatically manage resources
- Provide abstractions
- Allow interactive use
- Check for and report errors automatically
- Integrate tightly with `numpy`

PyCuda: Completeness

PyCuda exposes *all* of CUDA.

For example:


- Arrays and Textures
- Pagedlocked host memory
- Memory transfers (asynchronous, structured)
- Streams and Events
- Device queries
- (GL Interop)

PyCuda: Completeness

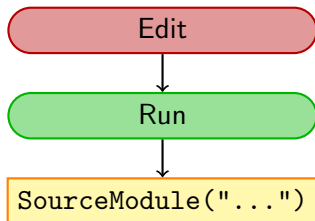
PyCuda supports every OS that CUDA supports.

- Linux
- Windows
- OS X


PyCuda: Workflow



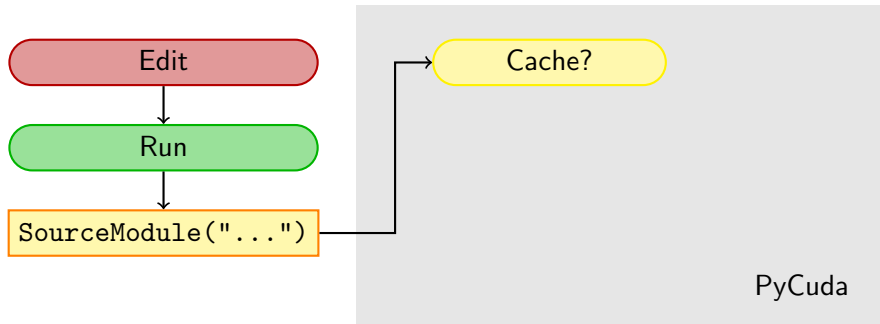
PyCuda: Workflow




PyCuda: Workflow




PyCuda: Workflow




PyCuda: Workflow



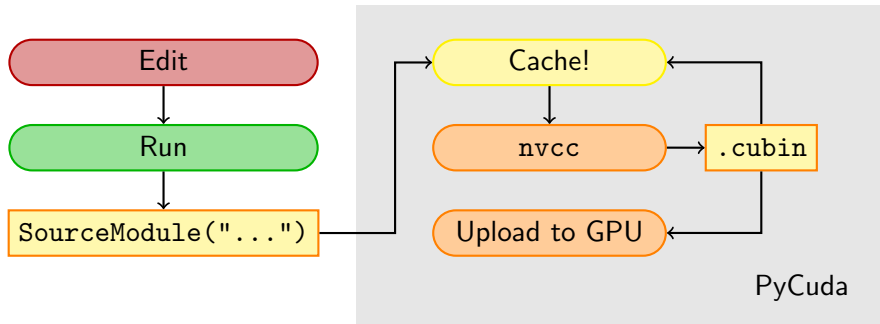
PyCuda: Workflow



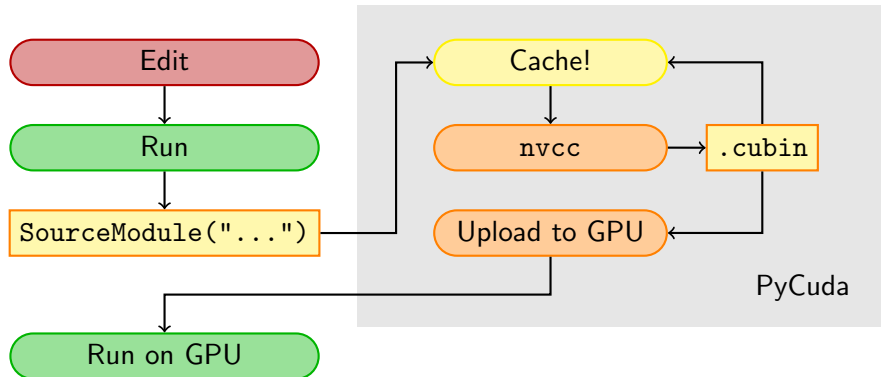
PyCuda: Workflow



PyCuda: Workflow



PyCuda: Workflow



Kernel Invocation: Automatic Copies


```
mod = pycuda.driver.SourceModule(  
    "__global__ my_func(float *out, float *in ){...} "  
func = mod.get_function("my_func")  
  
src = numpy.random.randn(400).astype(numpy.float32)  
dest = numpy.empty_like(src)
```

```
my_func(  
    cuda.Out(dest),  
    cuda.In(src),  
    block=(400,1,1))
```

- “InOut” exists, too.
- Only for immediate invocation style.

Automatic Cleanup


- Reachable objects (memory, streams, ...) are never destroyed.
- Once unreachable, released at an unspecified future time.
- Scarce resources (memory) can be explicitly freed. (`obj.free()`)
- Correctly deals with multiple contexts and dependencies.



gpuarray: Simple Linear Algebra

`pycuda.gpuarray`:

- Meant to look and feel just like `numpy`.
 - `gpuarray.to_gpu(numpy_array)`
 - `numpy_array = gpuarray.get()`
- No: `nd` indexing, slicing, etc. (yet!)
- Yes: `+`, `-`, `*`, `/`, `fill`, `sin`, `exp`, `rand`, `take`, ...
- Random numbers using `pycuda.curandom`
- Mixed types (`int32 + float32 = float64`)
- `print gpuarray` for debugging.
- Memory behind `gpuarray` available as `.gpudata` attribute.
 - Use as kernel arguments, textures, etc.



gpuarray: Elementwise expressions

Avoiding extra store-fetch cycles for elementwise math:

```
from pycuda.curandom import rand as curand
a_gpu = curand((50,))
b_gpu = curand((50,))

from pycuda.elementwise import ElementwiseKernel
lin_comb = ElementwiseKernel(
    "float a, float *x, float b, float *y, float *z",
    "z[i] = a*x[i] + b*y[i]")

c_gpu = gpuarray.empty_like(a_gpu)
lin_comb(5, a_gpu, 6, b_gpu, c_gpu)

assert la.norm((c_gpu - (5*a_gpu+6*b_gpu)).get()) < 1e-5
```

PyCuda: Vital Information

- <http://mathematician.de/software/pycuda>
- X Consortium License
(no warranty, free for all use)
- Requires: numpy, Boost C++,
Python 2.4+.
- Support via mailing list.



Questions?

?

Outline

1 Introduction


2 Programming GPUs

3 GPU Scripting

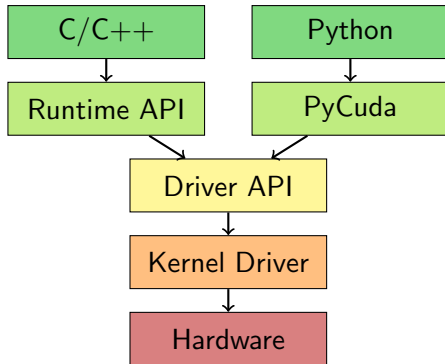
- Scripting + GPUs: A good combination
- Whetting your Appetite
- Working with PyCuda
- A peek under the hood
- Metaprogramming CUDA

4 PyCuda Hands-on: Matrix Multiplication

CUDA APIs



CUDA APIs



CUDA has two Programming Interfaces:

- “Runtime” high-level (`libcudart.so`, in the “toolkit”)
- “Driver” low-level (`libcuda.so`, comes with GPU driver)

(mutually exclusive)

Runtime vs. Driver API

Runtime \leftrightarrow Driver differences:

- Explicit initialization.

Runtime vs. Driver API

Runtime \leftrightarrow Driver differences:

- Explicit initialization.
- Code objects (“Modules”) become programming language objects.

Runtime vs. Driver API

Runtime \leftrightarrow Driver differences:

- Explicit initialization.
- Code objects (“Modules”) become programming language objects.
- Texture handling requires slightly more work.

Runtime vs. Driver API

Runtime \leftrightarrow Driver differences:

- Explicit initialization.
- Code objects (“Modules”) become programming language objects.
- Texture handling requires slightly more work.
- Only needs `nvcc` for compiling GPU code.

Runtime vs. Driver API

Runtime ↔ Driver differences:

- Explicit initialization.
- Code objects (“Modules”) become programming language objects.
- Texture handling requires slightly more work.
- Only needs `nvcc` for compiling GPU code.

Driver API:

- Conceptually cleaner
- Less sugar-coating (provide in Python)
- Not very different otherwise

PyCuda: API Tracing

With `./configure --cuda-trace=1`:

PyCuda: API Tracing

With `./configure --cuda-trace=1:`

```
import pycuda.driver as cuda
import pycuda.autoint
import numpy

a = numpy.random.randn(4,4).astype(numpy.float32)
a_gpu = cuda.mem_alloc(a.nbytes)
cuda.memcpy_htod(a_gpu, a)

mod = cuda.SourceModule("""
    __global__ void doublify( float *a)
    {
        int idx = threadIdx.x + threadIdx.y*4;
        a[idx] *= 2;
    }
    """)

func = mod.get_function(" doublify")
func(a_gpu, block=(4,4,1))

a_doubled = numpy.empty_like(a)
cuda.memcpy_dtoh(a_doubled, a_gpu)
print a_doubled
print a
```

```
cuInit
cuDeviceGetCount
cuDeviceGet
cuCtxCreate
cuMemAlloc
cuMemcpyHtoD
cuCtxGetDevice
cuDeviceComputeCapability
cuModuleLoadData
cuModuleGetFunction
cuFuncSetBlockShape
cuParamSetv
cuParamSetSize
cuLaunchGrid
cuMemcpyDtoH
cuCtxPopCurrent
cuCtxPushCurrent
cuMemFree
cuCtxPopCurrent
cuCtxPushCurrent
cuModuleUnload
cuCtxPopCurrent
cuCtxDestroy
```

Questions?

?

Outline

1 Introduction

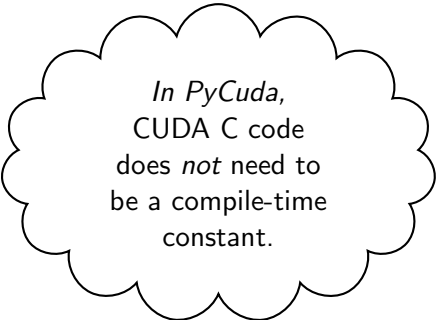
2 Programming GPUs

3 GPU Scripting

- Scripting + GPUs: A good combination
- Whetting your Appetite
- Working with PyCuda
- A peek under the hood
- **Metaprogramming CUDA**

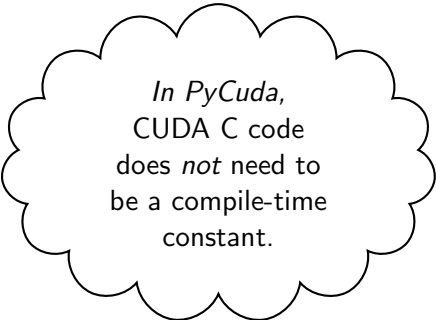
4 PyCuda Hands-on: Matrix Multiplication

Human vs Machine



In PyCuda,
CUDA C code
does *not* need to
be a compile-time
constant.

Human vs Machine

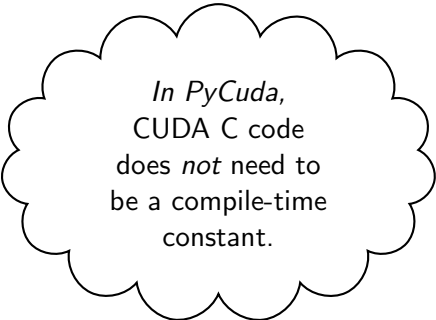


In PyCuda,
CUDA C code
does *not* need to
be a compile-time
constant.

(unlike the CUDA Runtime API)

Human vs Machine



Idea



In PyCuda,
CUDA C code
does *not* need to
be a compile-time
constant.


(unlike the CUDA Runtime API)

Human vs Machine




(unlike the CUDA Runtime API)

Human vs Machine





Machine




(unlike the CUDA Runtime API)

Human vs Machine



(unlike the CUDA Runtime API)

Human vs Machine



Easy to write

In PyCuda,
CUDA C code
does *not* need to
be a compile-time
constant.

(unlike the CUDA Runtime API)

Metaprogramming: machine-generated code

Why machine-generated code?



- Automated Tuning
(cf. ATLAS, FFTW)
- Data types
- Specialize code for given problem
- Constants faster than variables
(→ register pressure)
- Loop Unrolling

PyCuda: Support for Metaprogramming

- Access properties of compiled code:
`func.{num_regs,shared_size_bytes,local_size_bytes}`
- Exact GPU timing via events
- Can calculate hardware-dependent MP occupancy
- codepy (by Andreas):
 - Build C syntax trees from Python
 - Generates readable, indented C
- Or use a templating engine (many available, e.g. Cheetah)

Questions?

?

Outline

- 1 Introduction
- 2 Programming GPUs
- 3 GPU Scripting
- 4 PyCuda Hands-on: Matrix Multiplication**
 - Simple
 - Tiled
 - Meta-programming / Auto-tuning

Outline

- 1 Introduction
- 2 Programming GPUs
- 3 GPU Scripting
- 4 PyCuda Hands-on: Matrix Multiplication
 - Simple
 - Tiled
 - Meta-programming / Auto-tuning

Goal

- Multiply two (small) square matrices together.
- Use global memory only.
- Use a *single* block of threads.
- Each thread computes one element of the resulting matrix.

Instructions

- 1 `cd 3-pycuda-matrixmul-simple`
- 2 Edit `matrixmul_simple.py`
- 3 Complete the *TODOs*.

Code

Initialization:

```
import numpy as np
from pycuda import driver, compiler, gpuarray, tools
import atexit

# -- initialize the device
# the following lines are equivalent to "import pycuda.autotinit"
# only if GPU_NUMBER = 0
GPU_NUMBER = 0 # TODO: change me
driver . init ()
assert ( driver . Device . count () >= 1 )
dev = tools . get_default_device ( GPU_NUMBER )
ctx = dev . make_context ()
atexit . register ( ctx . pop )
```

Code

Memory allocation and transfer:

```
# define the (square) matrix size  
# note that we'll only use *one* block of threads here  
# as a consequence this number (squared) can't exceed max_threads,  
# see http://document.tician.de/pycuda/util.html#pycuda.tools.DeviceData  
# for more information on how to get this number for your device  
MATRIX_SIZE = 2  
  
# create two random square matrices  
a_cpu = np.random.randn(MATRIX_SIZE, MATRIX_SIZE).astype(np.float32)  
b_cpu = np.random.randn(MATRIX_SIZE, MATRIX_SIZE).astype(np.float32)  
  
# compute reference on the CPU to verify GPU computation  
c_cpu = np.dot(a_cpu, b_cpu)  
  
# transfer host (CPU) memory to device (GPU) memory  
a_gpu = gpuarray.to_gpu(a_cpu)  
b_gpu = gpuarray.to_gpu(b_cpu)  
  
# create empty gpu array for the result (C = A * B)  
c_gpu = gpuarray.empty((MATRIX_SIZE, MATRIX_SIZE), np.float32)
```

Code

GPU code compilation and execution:

```
# by specifying the constant MATRIX_SIZE
kernel_code = kernel_code_template % {
    'MATRIX_SIZE': MATRIX_SIZE
}

# compile the kernel code
mod = compiler.SourceModule(kernel_code)

# get the kernel function from the compiled module
matrixmul = mod.get_function("MatrixMulKernel")

# call the kernel on the card
matrixmul(
    # inputs
    a_gpu, b_gpu,
    # output
    c_gpu,
    # (only one) block of MATRIX_SIZE x MATRIX_SIZE threads
    block = (MATRIX_SIZE, MATRIX_SIZE, 1),
)
```


Code

GPU kernel code:

```
kernel_code_template = """
```

```
--global__ void MatrixMulKernel(float *a, float *b, float *c)
{
    // 2D Thread ID (assuming that only *one* block will be executed)
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue is used to store the element of the matrix
    // that is computed by the thread
    float Pvalue = 0;

    // Each thread loads one row of M and one column of N,
    // to produce one element of P.
    for (int k = 0; k < %(MATRIX.SIZE)s; ++k) {
        float Aelement = a[ ... ]; // TODO
        float Belement = b[ ... ]; // TODO
        Pvalue += Aelement * Belement;
    }

    // Write the matrix to device memory;
    // each thread writes one element
    c[ ... ] = Pvalue; // TODO
}
```

```
"""
```

Code

GPU kernel code (solution):

```
kernel_code_template = """
```

```
--global__ void MatrixMulKernel(float *a, float *b, float *c)
{
    // 2D Thread ID (assuming that only *one* block will be executed)
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue is used to store the element of the matrix
    // that is computed by the thread
    float Pvalue = 0;

    // Each thread loads one row of M and one column of N,
    // to produce one element of P.
    for (int k = 0; k < %(MATRIX.SIZE)s; ++k) {
        float Aelement = a[ty * %(MATRIX.SIZE)s + k];
        float Belement = b[k * %(MATRIX.SIZE)s + tx];
        Pvalue += Aelement * Belement;
    }

    // Write the matrix to device memory;
    // each thread writes one element
    c[ty * %(MATRIX.SIZE)s + tx] = Pvalue;
}
```

```
"""
```

Outline

- 1 Introduction
- 2 Programming GPUs
- 3 GPU Scripting
- 4 PyCuda Hands-on: Matrix Multiplication**
 - Simple
 - **Tiled**
 - Meta-programming / Auto-tuning

Goal

- Multiply two square matrices together.
- Use global memory and shared memory.
- Each thread block is assigned a "tile" of the resulting matrix and is responsible for generating the elements in that tile.
- Each thread in a block computes one element of the tile.

Code

GPU kernel code:

```
kernel_code_template = """
```

```
__global__ void MatrixMulKernel(float *A, float *B, float *C)
{

    const uint wA = %(MATRIX_SIZE)s;
    const uint wB = %(MATRIX_SIZE)s;

    // Block index
    const uint bx = blockIdx.x;
    const uint by = blockIdx.y;

    // Thread index
    const uint tx = threadIdx.x;
    const uint ty = threadIdx.y;

    // Index of the first sub-matrix of A processed by the block
    const uint aBegin = wA * %(BLOCK_SIZE)s * by;
    // Index of the last sub-matrix of A processed by the block
    const uint aEnd = aBegin + wA - 1;
    // Step size used to iterate through the sub-matrices of A
    const uint aStep = %(BLOCK_SIZE)s;

    // Index of the first sub-matrix of B processed by the block
    const uint bBegin = %(BLOCK_SIZE)s * bx;
    // Step size used to iterate through the sub-matrices of B
    const uint bStep = %(BLOCK_SIZE)s * wB;
```

Code

GPU kernel code (cont'd):

```
// The element of the block sub-matrix that is computed
// by the thread
float Csub = 0;
// Loop over all the sub-matrices of A and B required to
// compute the block sub-matrix
for (int a = aBegin, b = bBegin;
     a <= aEnd;
     a += aStep, b += bStep)
{
    // Shared memory for the sub-matrix of A
    __shared__ float As[(BLOCK_SIZE)s][(BLOCK_SIZE)s];
    // Shared memory for the sub-matrix of B
    __shared__ float Bs[(BLOCK_SIZE)s][(BLOCK_SIZE)s];

    // Load the matrices from global memory to shared memory;
    // each thread loads one element of each matrix
    As[ty][tx] = A[a + wA * ty + tx];
    Bs[ty][tx] = B[b + wB * ty + tx];
    // Synchronize to make sure the matrices are loaded
    __syncthreads ();
}
```

Code

GPU kernel code (cont'd):

```

// Multiply the two matrices together;
// each thread computes one element
// of the block sub-matrix
for (int k = 0; k < %(BLOCK_SIZE)s; ++k)
    Csub += As[ty][k] * Bs[k][tx];

// Synchronize to make sure that the preceding
// computation is done before loading two new
// sub-matrices of A and B in the next iteration
    __syncthreads ();
}

// Write the block sub-matrix to global memory;
// each thread writes one element
const uint c = wB * %(BLOCK_SIZE)s * by + %(BLOCK_SIZE)s * bx;
C[c + wB * ty + tx] = Csub;
}

```

"""

Outline

- 1 Introduction
- 2 Programming GPUs
- 3 GPU Scripting
- 4 PyCuda Hands-on: Matrix Multiplication**
 - Simple
 - Tiled
 - **Meta-programming / Auto-tuning**

Goal

- Multiply two matrices together (any size).
- Use global memory and shared memory.
- Implement various optimizations:
 - different granularities of parallelism (block and work sizes),
 - loop unrolling,
 - register pressure (spilling),
 - pre-fetching (global memory load).
- “Instrumentalize” the code using a *template engine* (Cheetah).
- Auto-tune depending on the hardware *and* the input data.

Instructions


- 1 `cd 5-pycuda-matrixmul-opt`
- 2 Implement your auto-tuning function.
- 3 Use PyCuda to gather informations (registers, occupancy).

Code

Show the code ;-)

Some numbers

3D Filterbank Convolutions used in our Visual Cortex Simulations:



Conclusions

- GPUs (or something like them) are here to stay.

Conclusions

- GPUs (or something like them) are here to stay.
- First factor of 5-10 is usually easy to reach.

Conclusions

- GPUs (or something like them) are here to stay.
- First factor of 5-10 is usually easy to reach.
- Next factor of 5-10 is a little bit harder.

Conclusions

- GPUs (or something like them) are here to stay.
- First factor of 5-10 is usually easy to reach.
- Next factor of 5-10 is a little bit harder.
- Next factor of 5-10 is a lot harder.
 - Requires deep understanding of the hardware architecture.
 - Usually involves significant rethinking of algorithm.

Conclusions

- GPUs (or something like them) are here to stay.
- First factor of 5-10 is usually easy to reach.
- Next factor of 5-10 is a little bit harder.
- Next factor of 5-10 is a lot harder.
 - Requires deep understanding of the hardware architecture.
 - Usually involves significant rethinking of algorithm.
- GPUs and scripting work surprisingly well together.
 - Favorable balance btw ease-of-use and raw performance.
 - Enable (easy) Metaprogramming.

Conclusions

- GPUs (or something like them) are here to stay.
- First factor of 5-10 is usually easy to reach.
- Next factor of 5-10 is a little bit harder.
- Next factor of 5-10 is a lot harder.
 - Requires deep understanding of the hardware architecture.
 - Usually involves significant rethinking of algorithm.
- GPUs and scripting work surprisingly well together.
 - Favorable balance btw ease-of-use and raw performance.
 - Enable (easy) Metaprogramming.
- Python / PyCuda rocks!

Thank you

